geoRglm : a package for generalised linear spatial models
introductory session

Ole F. Christensen & Paulo J. Ribeiro Jr.

Last update: June 30, 2010

The objective of this page is to introduce the reader to the **geoRglm** commands and show how they can be used. The commands used here are basic examples of the package handling, where we typically use default arguments for the function calls. We encourage the user also to inspect other function arguments.

For further details on the functions included in **geoRglm**, we refer to the **geoRglm** documentation.

1 STARTING A SESSION AND LOADING DATA

After starting an R session, we first load **geoR** and **geoRglm** with the commands:

```
> library(geoR)
> library(geoRglm)
```

If the installation directories for the packages are not the default locations for R packages, type:

```
  library(geoR, lib.loc = "PATH_TO_geoR")
  library(geoRglm, lib.loc = "PATH_TO_geoRglm")
```

where PATH_TO_geoR and PATH_TO_geoRglm are the paths to the directories where **geoR** and **geoRglm** are installed, respectively. If **geoRglm** is correctly loaded the following message will be displayed:

```
----------------------------------------------------------
geoRglm - a package for generalised linear spatial models
geoRglm version 0.8-28 (2010-30-06) is now loaded
----------------------------------------------------------
```

Helpfiles are available for **geoRglm**. For getting help on the function `glsm.mcmc`, just type:

```
> help(glsm.mcmc)
```

2 DATA AND MODELS

Spatial generalised linear mixed models are GLM's with spatially correlated random effects. The two interesting types of data are Poisson and Binomial data.

Typically, data are stored as an object (a list) of class `"geodata"` (see the **geoR** introductory session for more details on this). For the data sets considered here, the object will sometimes include a vector `units.m` consisting of, observation times or size of area for the Poisson distribution, or number of trials for the binomial distribution.

We use the data sets `b50` and `p50` included in the **geoRglm** distribution for the examples presented in this document. These data sets can be loaded by typing:

```
> data(b50)
> data(p50)
```

To gain an understanding of the Poisson model, we will here generate a simulation from such a models as follows. First we use the **geoR** function `grf` to generate a simulation from a Gaussian random field.

```
> sim.g <- grf(grid = expand.grid(x = seq(1, 10, l = 10),
+      y = seq(1, 10, l = 10)), cov.pars = c(0.1, 0.2))
```

Assume that we wish to simulate from a Poisson model with log-link, and assume that the observation times (or observation areas) vary between the locations. A simulation from the model is obtained by simulating from the Poisson distribution with the appropriate condtitional mean.

```
> sim <- list(coords = sim.g$coords, units.m = c(rep(1,
+      50), rep(5, 50)))
> attr(sim, "class") <- "geodata"
> sim$data <- rpois(100, lambda = sim$units.m * exp(sim.g$data))
```

Observe that the upper half of the figure corresponds to observation times equal to 5, where the simulated counts are larger.

```
> plot(sim$coords[, 1], sim$coords[, 2], type = "n")
> text(sim$coords[, 1], sim$coords[, 2], format(sim$data))
```
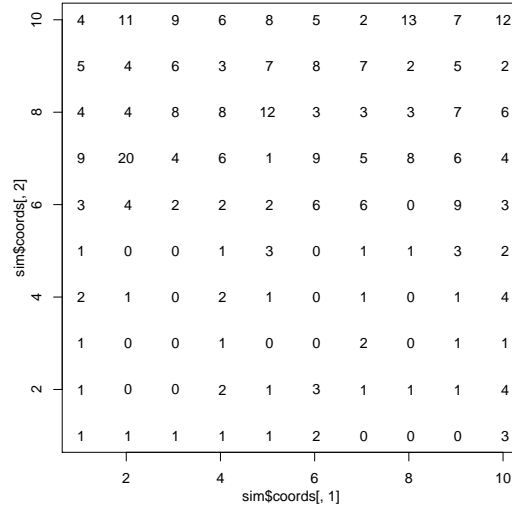

Figure 1: Simulated data

Exercise   Generate a simulation from a spatial binomial-logit model.

3   MCMC SIMULATION

The core part of **geoRglm** consist of generating MCMC simulations from the conditional distibution of the random effects at the data locations given the actual observed data. Such a simulation algorithm is needed for any likelihood inference in generalised linear spatial models (prediction, Bayesian inference and parameter estimation). Here we consider the fixed parameter case, which is implemented in the function `glsm.mcmc`.

The function uses a Langevin-Hastings MCMC algorithm for simulating from the conditional distribution. The nugget effect parameter (microscale variation) in the underlying Gaussian field can be set to a fixed value. The same applies for the smoothness and anisotropy parameters. Options for taking covariates (trends) into account are also included.

An example for Poisson data where we assume a logarithmic link, and where all parameters are fixed is shown below (for illustration purposes, some parameter values are just taken).

First we need to tune the algorithm by scaling the proposal variance so that acceptance rate is approximately 60 percent (optimal acceptance rate for Langevin-Hastings algorithm). This is done by trial and error.

3

```
> model2 <- list(cov.pars = c(1, 1), beta = 1, family = "poisson")
> mcmc2.test <- mcmc.control(S.scale = 0.2, thin = 1)
> test2.tune <- glsm.mcmc(p50, model = model2, mcmc.input = mcmc2.test)

iter. numb. 1000  : Acc.-rate =  0.877
MCMC performed: n.iter. =  1000 ; thinning =  1 ; burn.in =  0
```

After a few tryouts we decide to use `S.scale = 0.5`.

```
> mcmc2.tune <- mcmc.control(S.scale = 0.5, thin = 1)
> test2.tune <- glsm.mcmc(p50, model = model2, mcmc.input = mcmc2.tune)

iter. numb. 1000  : Acc.-rate =  0.57
MCMC performed: n.iter. =  1000 ; thinning =  1 ; burn.in =  0
```

We also need to study convergence of the chain and how well the chain is mixing. For this we use the functions in the **coda** package. We first load **coda** and then create a `mcmc` object to be used. Only a glimpse of the functionallity in **coda** is shown here, we encourage the reader to investigate further.

```
> library(coda)
> test2.tune.c <- create.mcmc.coda(test2.tune, mcmc.input = mcmc2.tune)
```

We encourage the user to always make traceplots and autocorrelation plots for all variables (in this case all 50 random effects). But in order to not to clutter this introduction with a large number of plots, we will present these plots for only one variable.

```
> test2.tune.c <- create.mcmc.coda(test2.tune$simulations[45,
+     ], mcmc.input = list(S.scale = 0.5, thin = 1))
> par(mfrow = c(1, 2))
> plot(test2.tune.c, density = FALSE, ask = FALSE, auto.layout = FALSE)
> autocorr.plot(test2.tune.c, ask = FALSE, auto.layout = FALSE)
```
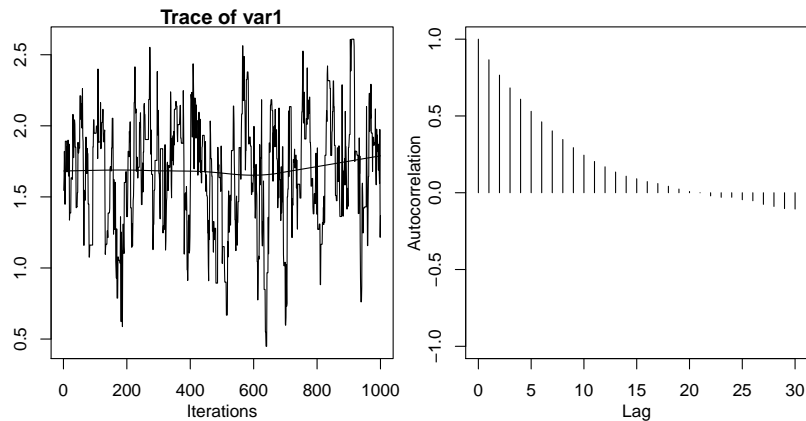

Figure 2: Results from `glsm.mcmc` for one variable.

To reduce the autocorrelation of the samples we decide to subsample every 10 iterations (default); when working with larger data sets we may need to make a more extensive subsampling, say, storing only every 100 iterations.

```
> mcmc2 <- mcmc.control(S.scale = 0.5)
> test2 <- glsm.mcmc(p50, model = model2, mcmc.input = mcmc2)

iter. numb. 1000   : Acc.-rate =  0.596
iter. numb. 2000   : Acc.-rate =  0.607
iter. numb. 3000   : Acc.-rate =  0.602
iter. numb. 4000   : Acc.-rate =  0.61
iter. numb. 5000   : Acc.-rate =  0.611
iter. numb. 6000   : Acc.-rate =  0.595
iter. numb. 7000   : Acc.-rate =  0.611
iter. numb. 8000   : Acc.-rate =  0.615
iter. numb. 9000   : Acc.-rate =  0.61
iter. numb. 10000  : Acc.-rate =  0.625
MCMC performed: n.iter. =  10000 ; thinning =  10 ; burn.in =  0
```

Exercise Produce traceplots and autocorrelation plots using the **coda** functions above (where you probably want to use default values of the arguments `ask` and `auto.layout`) for the MCMC output contained in `test2`.

5

4 SPATIAL PREDICTION

For the model and data above we now consider spatial prediction, assuming that parameters are fixed. Full Bayesian prediction methods are also implemented and will be presented in Section 5.

For computational reasons we consider prediction at only two locations here. Minimal mean square error prediction of the intensity at the two locations (0.5, 0.5) and (1, 0.4). Here we use the object `test2` created in the previous section

```
> out2 <- output.glm.control(sim.predict = TRUE)
> pred.test2 <- glsm.krige(test2, locations = cbind(c(0.5,
+     0.5), c(1, 0.4)), output = out2)

 glsm.krige: Prediction for a generalised linear spatial model
```

The output is a list including the predicted values (`pred.test2$predict`), the prediction variances (`pred.test2$krige.var`) and the estimated Monte Carlo standard errors on the predicted values (`pred.test2$mcmc.error`). Printing out the predicted values and the associated Monte Carlo standard errors:

```
> cbind(pred.test2$predict, pred.test2$mcmc.error)

          [,1]       [,2]
[1,] 5.503865 0.04906689
[2,] 4.934100 0.04225835
```

we see that the Monte Carlo standard errors (the errors due to the MCMC-simulation) are small compared to predicted values, which is very satisfactory.

By specifying `sim.predict = TRUE`, simulations are drawn from the predictive intensity at the two prediction locations (`pred.test2$simulations`). These simulations are plotted in Figure 3.

Exercise Make prediction on a 40 by 40 regular grid, and afterwards visualise predictions using the **geoR** function `image`. We suggest to consult the section about prediction in the **geoR** introduction before starting.

Exercise Binomial data can be specified by `family="binomial"` in `glsm.mcmc`. The exercise consist in repeating the commands above for the binomial data set `b50`.
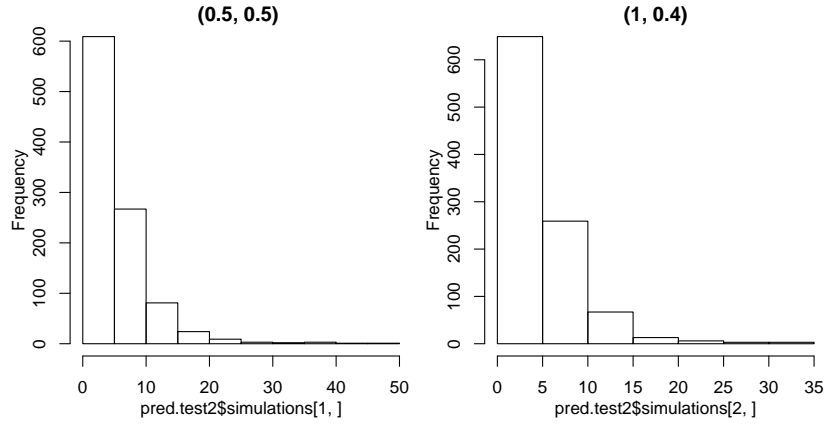
Figure 3: Histograms for simulations from the predictive distribution for two locations.

Remark In the classical Gaussian model there exist different kriging flavours : simple kriging, ordinary kriging, etc. The prediction above for a GLSM corresponds to simple kriging because all parameters are fixed. Ordinary and Universal kriging in the Gaussian model corresponds to using a uniform flat prior on the beta parameter. For a GLSM, the corresponding prediction method is not implemented in the functions above. However, a uniform flat prior on β parameter is considered in the next section on Bayesian inference, and in addition it is also implemented for fixed covariance parameters in the two functions `pois.krige` and `binom.krige`.

5 BAYESIAN ANALYSIS

Bayesian analysis for the Poisson-log normal model and the binomial-logit model is implemented by the functions `pois.krige.bayes` and `binom.krige.bayes`, respectively. Model parameters can be treated as fixed or random.

As an example consider first a model without nugget and including uncertainty in the β and σ^2 parameters (mean and variance of the random effects S, respectively). A Bayesian analysis is made by typing commands like:

```
> prior5 <- prior.glm.control(phi.prior = "fixed", phi = 0.1)
> mcmc5.tune <- mcmc.control(S.scale = 0.01, thin = 1)
> test5.tune <- pois.krige.bayes(p50, prior = prior5, mcmc.input = mcmc5.tune)

pois.krige.bayes: model with mean being constant
iter. numb. 1000 ; Acc.-rate = 0.98
```

```
MCMC performed: n.iter. =  1000 ; thinning =  1 ; burn.in =  0
Only Bayesian estimation of model parameters
```

Now chose `S.scale` (Acc-rate=0.60 is preferable). After having adjusted the parameters for the MCMC algorithm and checking the output we run an analysis (where we here omit the printing of the messages from the MCMC iterations for brevity).

```
> mcmc5 <- mcmc.control(S.scale = 0.075, thin = 100)
> out5 <- output.glm.control(threshold = 10, quantile = c(0.05,
+     0.99))
> test5 <- pois.krige.bayes(p50, locations = t(cbind(c(2.5,
+     3), c(-6050, -3270))), prior = prior5, mcmc.input = mcmc5,
+     output = out5)
```

The output is a list which contains the five arguments `posterior`, `predictive`, `model`, `prior` and `mcmc.input`. The `posterior` contains information on the posterior distribution of the parameters, and the conditional simulations of the signal $g^{-1}(S)$ at the data locations. The `predictive` contains information on the predictions, where `predictive$median` is the predicted signal and `predictive$uncertainty` is the associated uncertainty. The `threshold = 10` argument gives probabilities of the predictive distribution of the signal being less than 10 (`test5$predictive$probability`). The `quantiles = c(0.05,0.99)` gives the 0.05 and 0.99 quantiles of the predictive distribution of the signal (`test5$predictive$quantiles`).

Below we show the simulations from the posterior distribution of the signal at a few data locations.

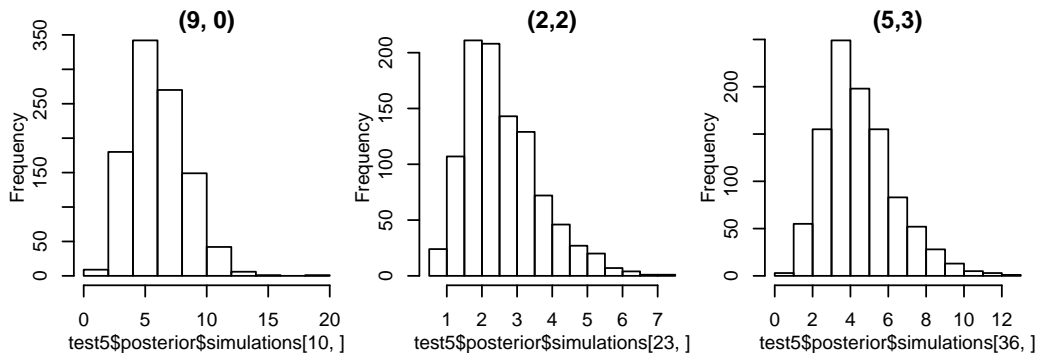
Figure 4: Histograms

Now we consider an example with a random correlation scale parameter phi and a positive nugget for the random effects S. The program is using a

discretised prior for phi, where the discretisation is given by the argument `phi.discrete`). The argument `tausq.rel = 0.05` gives the relative nugget for S, i.e. the relative microscale variation.

```
> mcmc6.tune <- mcmc.control(S.scale = 0.075, n.iter = 2000,
+     thin = 100, phi.scale = 0.01)
> prior6 <- prior.glm.control(phi.prior = "uniform", phi.discrete = seq(0.02,
+     1, 0.02), tausq.rel = 0.05)
> test6.tune <- pois.krige.bayes(p50, prior = prior6, mcmc.input = mcmc6.tune)
```

Acc-rate=0.60 , acc-rate-phi = 0.25-0.30 are preferable. After having adjusted the parameters for the MCMC algorithm and checking the output we run an analysis.

WARNING: RUNNING THE NEXT COMMAND CAN BE TIME-CONSUMING

```
> mcmc6 <- mcmc.control(S.scale = 0.075, n.iter = 4e+05,
+     thin = 200, burn.in = 5000, phi.scale = 0.12, phi.start = 0.5)
> test6 <- pois.krige.bayes(p50, locations = t(cbind(c(2.5,
+     3.5), c(-60, -37))), prior = prior6, mcmc.input = mcmc6)
```

Below we show the posterior distribution of the two covariance parameters and the beta parameter.
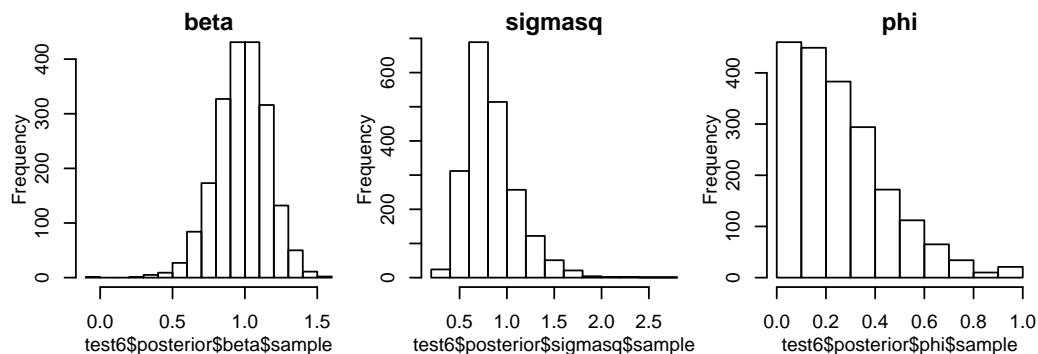

Figure 5: Samples from the posterior

Exercise Use **coda** and the function `create.mcmc.coda` to investigate the convergence and mixing of the MCMC algorithm for the examples above.

Exercise Construct similar commands as above using the function `binom.krige.bayes` on the data set `b50` yourself.

9

Remark The Bayesian inferential functions differ from the functions used in the fixed parameter case in the following two ways : A. there exist two functions, one for binomial data and for Poisson data. B. the Bayesian inferential procedure has not been split into two functions, a MCMC function and prediction function, similar to `glsm.mcmc` and `glsm.krige`. The main reason for these differences is historical, the functions `glsm.mcmc` and `glsm.krige` were introduced in geoRglm version 0.8-0 in the spring 2004. A reason for not restructuring the Bayesian functions similarly has been compatibility with the **geoR** function `krige.bayes` (in addition to the lack of time for doing so !).

6 ADDITIONAL INFORMATION

Package **geoRglm** also contain some functions for likelihood inference (MCMC-MLE). They are relatively slow to use, and have therefore not been included in this introduction.

We strongly encourage the user to study the relevant literature and also the **geoRglm** homepage before starting using the package.